

REPORT 3

Introduction to CUDA and OpenCL

AGH-Faculty of Physics and Applied Computer Science

authors:

Kinga Pyrek
Aleksandra Rolka

On our last lab classes we conducted experiment to learn more about the behavior of `cudaMallocManaged()` and we worked with The NVIDIA Visual Profiler. The Visual Profiler is a graphical profiling tool that displays a timeline of our application's CPU and GPU activity. We used it to show differences in the way of accessing UM (unified memory).

We also looked into a page fault, which is a type of exception raised by computer hardware. It occurs when a running program accesses a memory page that is not currently mapped by the memory management unit into the virtual address space of a process.


It means that the page may be accessible to the process, but requires a mapping to be added to the process page tables, and may additionally require the actual page contents to be loaded from a backing store such as a disk.

When handling a page fault, the operating system generally tries to make the required page accessible at the location physical memory, or terminates the program in case of an illegal memory access.

Contrary to what "fault" might suggest, valid page faults are not errors, and are common and necessary to increase the amount of memory


available to programs in any operating system that utilizes virtual memory.

There are results of 4 versions of the program:




Visual profiling results when unified memory is accessed only by the GPU

Access to unified memory has only GPU, that's why we have only GPU page faults. It doesn't take too much time.



Visual profiling results when unified memory is accessed only by the CPU


This time only CPU has access to UM and there are only CPU page faults, but its time is much longer and it performs right after `cudaMallocManaged` is done.



Visual profiling results when unified memory is accessed first by the GPU then the CPU

In third case we gave access to UM first to GPU and then to CPU. What's new, that Data Migration(DtoH) appeared. It shows which one has

first access and that GPU's and CPU's page fault are processed at the same time.



Visual profiling results when unified memory is accessed first by the CPU then the GPU

In last one we have reverse order access to unified memory and in this case GPU's page fault are processed not at the same time, but after CPU's page faults are done. We have shortest time of Data Migration, but overall time is longer.

Next we were analyzing files:

- vector_add_standard
- vector_add_prefetch_GPU
- vector_add_prefetch_GPU_init_GPU
- vector_add_prefetch_GPUCPU_init_GPU

to familiarize ourselves with data prefetching.

To **prefetch** data it means to fetch data in advance. The benefit of prefetching data is to leverage the asynchronous aspect of memory accesses in CUDA. When a memory access operation is executed, it does not block other operations following it as long as they don't use the data from the operation.

As written in (Kernel without prefetching), every addition waits for its data to be loaded from memory. Inside the loop of (Kernel with prefetching), the device first launches a memory load operation for the next iteration and does an addition in parallel. The time for the addition is actually

overlapping with the memory access time. But the increased register usage may lower the number of active warps on an SM.

These are the profiling results we've got:

```
[cuda-s13@lhcbgpu1 lab4]$ nvprof ./a.out
==10114== NVPROF is profiling process 10114, command: ./a.out
Success! All values calculated correctly.
==10114== Profiling application: ./a.out
==10114== Profiling result:
   Type  Time(%)   Time    Calls   Avg      Min      Max  Name
GPU activities: 100.00% 139.02ms    1 139.02ms 139.02ms 139.02ms addVectorsInto(float*, float*, float*, int)
API calls:      67.46% 339.59ms    3 113.20ms 24.584us 339.50ms cudaMallocManaged
              27.63% 139.07ms    1 139.07ms 139.07ms 139.07ms cudaDeviceSynchronize
              4.72% 23.782ms    3 7.9274ms 7.2425ms 8.3038ms cudaFree
              0.09% 457.60us    1 457.60us 457.60us 457.60us cuDeviceTotalMem
              0.06% 290.47us   96 3.0250us 838ns 93.308us cuDeviceGetAttribute
              0.02% 116.22us    1 116.22us 116.22us 116.22us cudaLaunchKernel
              0.01% 38.832us    1 38.832us 38.832us 38.832us cuDeviceGetName
              0.00% 16.273us    1 16.273us 16.273us 16.273us cudaGetDevice
              0.00% 10.337us    1 10.337us 10.337us 10.337us cuDeviceGetPCIBusId
              0.00% 3.7020us    3 1.2340us 978ns 1.6060us cuDeviceGetCount
              0.00% 3.7020us    1 3.7020us 3.7020us 3.7020us cuDeviceGetAttribute
              0.00% 2.5140us    2 1.2570us 908ns 1.6060us cuDeviceGet
              0.00% 1.1180us    1 1.1180us 1.1180us 1.1180us cuDeviceGetUuid
              0.00% 908ns      1 908ns    908ns 908ns  cudaGetLastError

==10114== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
   13192  27.974KB  4.0000KB 192.00KB  360.3906MB  41.45622ms  Host To Device
    766   171.03KB 4.0000KB 0.9961MB  127.9375MB  10.71466ms  Device To Host
    1013      -      -      -      -  125.6237ms  Gpu page fault groups
Total CPU Page faults: 1536
```

vector_add_standard's profiling results

This is the simplest version of our program. As we can see there are a lot of page faults(in GPU activity also).

```
[cuda-s13@lhcbgpu1 lab4]$ nvcc vector_add_prefetch_gpu.cu
[cuda-s13@lhcbgpu1 lab4]$ nvprof ./a.out
==10448== NVPROF is profiling process 10448, command: ./a.out
Success! All values calculated correctly.
==10448== Profiling application: ./a.out
==10448== Profiling result:
   Type  Time(%)   Time    Calls   Avg      Min      Max  Name
GPU activities: 100.00% 2.5613ms    1 2.5613ms 2.5613ms 2.5613ms addVectorsInto(float*, float*, float*, int)
API calls:      76.92% 302.50ms    3 100.83ms 21.441us 302.43ms cudaMallocManaged
              11.71% 46.032ms    1 46.032ms 46.032ms 46.032ms cudaDeviceSynchronize
              5.40% 21.253ms    3 7.0843ms 6.3248ms 7.6994ms cudaFree
              3.28% 12.912ms    3 4.3042ms 22.139us 12.713ms cudaMemPrefetchAsync
              2.50% 9.8447ms    1 9.8447ms 9.8447ms 9.8447ms cuDeviceTotalMem
              0.11% 413.53us   96 4.3070us 839ns 150.23us cuDeviceGetAttribute
              0.04% 173.42us    1 173.42us 173.42us 173.42us cuDeviceGetName
              0.02% 90.444us    1 90.444us 90.444us 90.444us cudaLaunchKernel
              0.01% 23.187us    1 23.187us 23.187us 23.187us cudaGetDevice
              0.00% 10.407us    1 10.407us 10.407us 10.407us cuDeviceGetPCIBusId
              0.00% 4.8190us    1 4.8190us 4.8190us 4.8190us cuDeviceGetAttribute
              0.00% 3.7010us    3 1.2330us 838ns 1.8160us cuDeviceGetCount
              0.00% 2.5130us    2 1.2560us 977ns 1.5360us cuDeviceGet
              0.00% 1.4670us    1 1.4670us 1.4670us 1.4670us cuDeviceGetUuid
              0.00% 1.0480us    1 1.0480us 1.0480us 1.0480us cudaGetLastError

==10448== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    192  2.0000MB  2.0000MB  2.0000MB  384.0000MB  34.35018ms  Host To Device
    768  170.67KB  4.0000KB  0.9961MB  128.0000MB  10.71642ms  Device To Host
Total CPU Page faults: 1536
```

vector_add_prefetch_GPU's profiling results

Here was used `cudaMemPrefetchAsync` function, which prefetches memory to the specified destination device. As a result of using that we do not have page faults in GPU, but we still have the same number of CPU page faults as in the previous sample.

```
[cuda-s13@lhcbgpu1 lab4]$ nvcc vector_add_prefetch_gpu_init_gpu.cu
[cuda-s13@lhcbgpu1 lab4]$ nvprof ./a.out
==10686== NVPROF is profiling process 10686, command: ./a.out
Success! All values calculated correctly.
==10686== Profiling application: ./a.out
==10686== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:  50.78%    2.5634ms      1    2.5634ms    2.5634ms    2.5634ms    addVectorsInto(float*, float*, float*, int)
                49.22%    2.4851ms      3    828.36us    827.78us    828.77us    initWith(float, float*, int)
API calls:      90.45%    297.49ms      3    99.165ms    22.908us    297.42ms    cudaMallocManaged
                6.77%    22.253ms      3    7.4177ms    4.7918ms    12.629ms    cudaFree
                1.53%    5.0187ms      1    5.0187ms    5.0187ms    5.0187ms    cudaDeviceSynchronize
                0.88%    2.8902ms      3    963.41us    939.16us    1.0037ms    cudaMemPrefetchAsync
                0.14%    460.53us      1    460.53us    460.53us    460.53us    cuDeviceTotalMem
                0.10%    322.46us      1    322.46us    322.46us    322.46us    cuDeviceGetName
                0.09%    308.56us     96    3.2140us      838ns    104.27us    cuDeviceGetAttribute
                0.03%    104.00us      4    25.999us     9.0800us    72.915us    cudaLaunchKernel
                0.00%    15.784us      1    15.784us    15.784us    15.784us    cudaGetDevice
                0.00%    10.058us      1    10.058us    10.058us    10.058us    cuDeviceGetPCIBusId
                0.00%    4.9590us      3    1.6530us     1.0480us    2.7940us    cuDeviceGetCount
                0.00%    2.9340us      2    1.4670us      978ns    1.9560us    cuDeviceGet
                0.00%    1.6060us      1    1.6060us    1.6060us    1.6060us    cudaDeviceGetAttribute
                0.00%    1.2570us      1    1.2570us    1.2570us    1.2570us    cuDeviceGetUuid
                0.00%      977ns       1      977ns      977ns      977ns    cudaGetLastError

==10686== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    768   170.67KB  4.0000KB  0.9961MB  128.0000MB  10.70000ms  Device To Host
Total CPU Page faults: 384
```

vector_add_prefetch_GPU_init_GPU's profiling results

In the third code there is `initWith` function, which initializes space for data on GPU, due to that the Host to Device transfer is gone.

Even though it is slower because of initializing space on GPU it repays the execution time raise because we save the memory. Also total amount of page faults is only 384, so it's 5 times smaller than in the previous cases.

```
[cuda-s13@lhcbgpu1 lab4]$ nvcc vector_add_prefetch_gpucpu_init_gpu.cu
[cuda-s13@lhcbgpu1 lab4]$ nvprof ./a.out
==10848== NVPROF is profiling process 10848, command: ./a.out
Success! All values calculated correctly.
==10848== Profiling application: ./a.out
==10848== Profiling result:
   Type  Time(%)    Time      Calls      Avg      Min      Max  Name
GPU activities: 50.78%  2.5683ms      1  2.5683ms  2.5683ms  2.5683ms  addVectorsInto(float*, float*, float*, int)
               49.22%  2.4890ms      3  829.67us  825.41us  832.65us  initWith(float, float*, int)
   API calls: 79.73%  310.55ms      3  103.52ms  23.536us  310.48ms  cudaMallocManaged
               11.95%  46.548ms      4  11.637ms  943.49us  43.648ms  cudaMemPrefetchAsync
               6.76%  26.314ms      3  8.7713ms  4.7677ms  16.639ms  cudaFree
               1.29%  5.0256ms      1  5.0256ms  5.0256ms  5.0256ms  cudaDeviceSynchronize
               0.12%  473.25us      1  473.25us  473.25us  473.25us  cuDeviceTotalMem
               0.11%  432.81us     96  4.5080us    838ns  220.00us  cuDeviceGetAttribute
               0.03%  101.62us      4  25.404us  9.2190us  70.120us  cudaLaunchKernel
               0.01%  41.975us      1  41.975us  41.975us  41.975us  cuDeviceGetName
               0.00%  18.438us      1  18.438us  18.438us  18.438us  cudaGetDevice
               0.00%  10.826us      1  10.826us  10.826us  10.826us  cuDeviceGetPCIBusId
               0.00%  3.6310us      3  1.2100us    908ns  1.7460us  cuDeviceGetCount
               0.00%  2.6540us      2  1.3270us    978ns  1.6760us  cuDeviceGet
               0.00%  1.6760us      1  1.6760us  1.6760us  1.6760us  cudaDeviceGetAttribute
               0.00%  1.1880us      1  1.1880us  1.1880us  1.1880us  cuDeviceGetUuid
               0.00%    978ns      1    978ns    978ns    978ns  cudaGetLastError

==10848== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
      64  2.0000MB  2.0000MB  2.0000MB  128.0000MB  10.2160ms  Device To Host
```

vector_add_prefetch_GPUCPU_init_GPU's profiling results

The last version of code is the most efficient one. The memory is initialized on CPU and GPU and we use prefetching data. In the wake of refinements there is no page faults.